Multi-layer Stack Ensembles for Time Series Forecasting

Nathanael Bosch[1,*] Oleksandr Shchur[2] Nick Erickson[2] Michael Bohlke-Schneider[2]
Caner Türkmen[2]

[1]Tübingen AI Center, University of Tübingen
[2]AWS
[*]Work done during an internship at AWS

Abstract   Ensembling is a powerful technique for improving the accuracy of machine learning models, with methods like stacking achieving strong results in tabular tasks. In time series forecasting, however, ensemble methods remain underutilized, with simple linear combinations still considered state-of-the-art. In this paper, we systematically explore ensembling strategies for time series forecasting. We evaluate 33 ensemble models—both existing and novel—across 50 real-world datasets. Our results show that stacking consistently improves accuracy, though no single stacker performs best across all tasks. To address this, we propose a multi-layer stacking framework for time series forecasting, an approach that combines the strengths of different stacker models. We demonstrate that this method consistently provides superior accuracy across diverse forecasting scenarios. Our findings highlight the potential of stacking-based methods to improve AutoML systems for time series forecasting.

1 Introduction

Time series forecasting plays a critical role in applications ranging from inventory management (Croston, 1972) to energy systems (Hong et al., 2016) and public health (Tsang et al., 2024). Driven by growing interest in the machine learning (ML) community, numerous forecasting models have been proposed in recent years. Yet no single model is universally dominant. Large-scale benchmarks show that while some models perform better on average, none consistently outperforms all others across diverse datasets and tasks (Aksu et al., 2024; Godahewa et al., 2021).

This variability in model performance motivates the use of automated machine learning (AutoML) for time series forecasting. By evaluating multiple models on a given dataset, an AutoML system can select the one that performs best for the specific forecasting task at hand. However, relying solely on model selection is often insufficient to achieve the lowest possible forecast error.

A growing body of work—both in academic research and prediction competitions—demonstrates that combining forecasts from multiple models can lead to substantial improvements in predictive performance (Bojer and Meldgaard, 2021; Makridakis et al., 2018, 2022). Yet, several key questions remain unanswered. While many forecast combination methods have been proposed over the years (X. Wang et al., 2023), there is limited understanding of which ensembling strategies are most effective in practice. In the related field of tabular AutoML, techniques such as multi-layer stacking repeatedly demonstrate state-of-the-art performance (Gijsbers et al., 2024). In contrast, time series forecasting has seen limited adoption of such techniques, with most practical systems relying on simple linear combinations of model outputs (Herzen et al., 2022; Shchur et al., 2023). Moreover, several studies suggest that in forecasting, simple arithmetic averages outperform more sophisticated ensembling techniques, which is often referred to as the "forecast combination puzzle" (Smith and Wallis, 2009; Stock and Watson, 2004).

To address this gap, we conduct a large-scale empirical study of forecast combination methods, aiming to identify strategies with the best predictive performance. Our main contributions are:

- **Empirical evaluation of ensembling methods**. We conduct a large-scale benchmark of 33 forecast combination methods across 50 real-world datasets on point and probabilistic forecasting tasks. Our results show that ensembling consistently improves the predictive accuracy, with learning-based ensembles significantly outperforming simple aggregation strategies.

- **Multi-layer stacking framework**. We propose a framework for building multi-layer stack ensembles for time series forecasting, designed to combine the strengths of different ensembling methods. This approach consistently outperforms any individual ensemble across both point and probabilistic forecasting tasks. Additional ablation studies reveal when and why multi-layer stacking is effective, demonstrating its adaptability and robustness across datasets.

2 Background

2.1 Probabilistic time series forecasting

Given a collection of N time series $\mathcal{D} = \{y_{i,1:T}\}_{i=1}^{N}$, with $y_{i,t} \in \mathbb{R}$, the goal of time series forecasting is to predict the future H values $y_{i,T+1:T+H}$ for each time series $i = 1, \ldots, N$. Specifically, in *probabilistic* time series forecasting, the aim is to model the conditional distribution $p(y_{i,T+1:T+H} \mid y_{i,1:T})$ for all i. Often, it is sufficient to produce a quantile forecast rather than the full distribution. Given a predefined set of quantile levels $\mathcal{Q} \subset (0, 1)$ with $|\mathcal{Q}| = Q$, the goal is to predict $\hat{y}_{i,T+h}^{q} \in \mathbb{R}$ such that $P(y_{i,T_i+h} < \hat{y}_{i,T_i+h}^{q}) = q$, for all q, i, h. Alternatively, if uncertainty quantification is not needed, a conditional mean or median forecast may suffice. In this work, we primarily focus on quantile forecasts, but all methods also apply to point forecasts and we will also evaluate these in Sec. 6.

Time series forecasting *models* $f_i : \mathbb{R}^T \rightarrow \mathbb{R}^{H \times Q}$ produce these quantile predictions, mapping $f_i : y_{i,1:T} \mapsto \hat{y}_{i,T+h}^{q}$ for all i. Here, the collection of models $\{f_i\}$ can be individually learned for each time series, an approach known as *local* time series modeling (Januschowski et al., 2020). Local time series models include *seasonal naive*, *exponential smoothing*, and *ARIMA*, which typically fit a small set of model parameters on each item individually, in order to capture simple patterns in the data such as trends and seasonality (R. Hyndman and Athanasopoulos, 2021).

On the other hand, a *global* time series model uses a single model for all time series, that is $f_i = f$ for all i. The function f is learned over the entire training set $\mathcal{D}$. The aim in global time series modeling is to use a higher capacity function f (e.g., a neural network) to capture both simple temporal patterns and any dynamics common to all time series in a dataset. This approach includes many recent deep learning-based approaches, such as *DeepAR* (Salinas, Flunkert, et al., 2020), *TFT* (Lim et al., 2021), *PatchTST* (Nie et al., 2022), *DLinear* (Zeng et al., 2023), or *TiDE* (Das et al., 2023).

Finally, *pretrained* models also fit parameterized functions to data, but are trained on large corpora of real-world and/or synthetic time series from diverse domains. The model is then applied as a zero-shot forecaster, that is, without any additional training on the dataset $\mathcal{D}$ associated with the given forecasting task. Recent examples of such pretrained models include *Chronos* (Ansari et al., 2024), *TimesFM* (Das et al., 2024), *Moirai* (Woo et al., 2024), and *TTM* (Ekambaram et al., 2024).

Time series models are trained via minimizing loss function $\mathcal{L} : \mathbb{R}^{H \times Q} \times \mathbb{R}^{H} \rightarrow \mathbb{R}_{\geq 0}$ over the training set. Examples include quantile loss for probabilistic forecasting, or mean absolute error for point forecasting (R. Hyndman and Athanasopoulos, 2021).

2.2 Forecast combination

While many approaches to modeling exist, none of these dominate others in terms of accuracy or efficiency. The best forecasting approach is therefore often a combination of different models (X. Wang et al., 2023). Let $f_{i,m}$ denote the forecast model for item i, obtained by using the modeling approach $m \in 1, \cdots, M$. Forecasts from these models can be combined with a combination function $g_i : (\mathbb{R}^{H \times Q})^M \rightarrow \mathbb{R}^{H \times Q}$ that operates on the outputs of a collection of base forecasting models
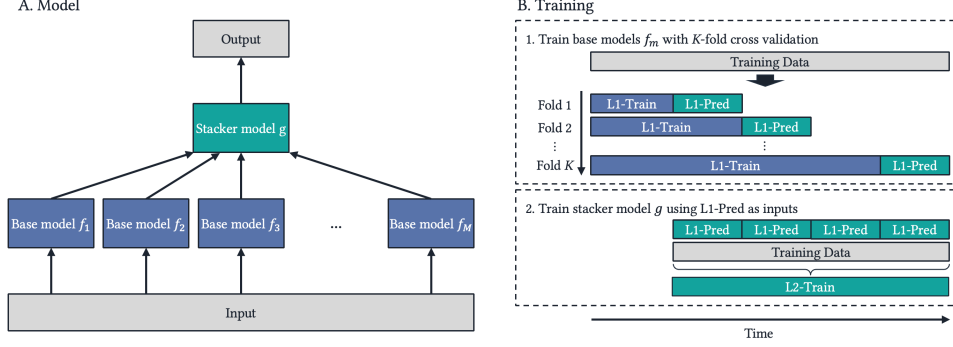
Figure 1: Architecture and training procedure of a single-layer stacker model.

$\{f_{i,m}\}$ to produce a single aggregated forecast. Composing the combination functions g_i with base forecasters $\{f_{i,m}\}$, we obtain forecast ensembles $\overline{f}_i$

$$\overline{f}_i : \mathbb{R}^T \to \mathbb{R}^{H \times Q}, \quad y_{i,1:T} \mapsto g_i(f_{i,1}(y_{i,1:T}), f_{i,2}(y_{i,1:T}), \ldots, f_{i,M}(y_{i,1:T})), \tag{1}$$

for all i. See also Fig. 1 for a visual description of the approach.

Similar to base models, combination methods can also be *global* with $g_i = g$ for all i. This is possible regardless of whether the base models are local or global. For example, we could learn global combinations g of individually fitted local models $\{f_{i,m}\}$. Conversely, we could use a different combination method g_i for each time series, when combining global forecasting methods $\{f_m\}$. Combinations of global and local models as base models are also possible. In the following, we will rely on this generality to drop the item index i from our notation. For example, in the following $\hat{y}_m$ may denote either the predictions for a single item $\hat{y}_{i,m}$ or for all items $\{y_{i,m}\}_{i=1}^N$, depending on the context.

Likely the most popular forecast combination approach is **simple averaging** via the mean or median of individual forecasts, e.g., $g(\hat{y}_1, \cdots, \hat{y}_M) = \frac{1}{M} \sum_m \hat{y}_m$. Simple averaging has been shown to often outperform individual forecasting models (Perrone and Cooper, 1995; X. Wang et al., 2023). **Model selection** can also be cast in the same notation. That is, $g(\hat{y}_1, \ldots, \hat{y}_M) = \hat{y}_{m^*}$, where m^* is the index of the model with the best validation score, i.e. $m^* \coloneqq \arg\min_{m=1,\ldots,M} \mathcal{L}(\hat{y}_m, y_{\text{out}})$, where y_{out} denotes the ground truth future values corresponding to the predictions $\hat{y}_m$.

In the AutoML literature, methods that combine outputs of individual models are referred to as *post-hoc ensembles* (Purucker and Beel, 2023) to differentiate them from other ensembling techniques like bagging and boosting. In forecasting literature, apart from *forecast combinations* (Clemen, 1989), these methods are called forecast ensembles (Adhikari and Agrawal, 2012) and model averaging (Montero-Manso et al., 2020). We use these terms interchangeably throughout the paper.

3 Stacking for time series forecasting

3.1 Stacking

Simple averaging and model selection are strong baseline approaches. In this work, however, we explore a broader class of methods that can *learn* the combination functions g from data. In ML literature, such methods that learn to combine the outputs of other models are often called *stacking* methods (Zhou, 2012, Ch. 4). We now review some model families that fall into this category.

Model averaging. One generalization of simple averaging is weighted model averaging, with a stacker model of the form $g(\hat{y}_1, \ldots, \hat{y}_M) = \sum_{m=1}^{M} \omega_m \hat{y}_m$, with weights $\omega \in \mathbb{R}^M$. The weights ω can be chosen based on the **model performance**. For example, the weights can be computed as $\omega_m \propto h(\mathcal{L}(\hat{y}_m, y_{\text{out}}))$, where h are chosen as $h_{\text{inv}}(L) = 1/L$, $h_{\text{sqr}}(L) = 1/L^2$, or $h_{\text{exp}}(L) = \exp(1/L)$, subject to $\sum_{m=1}^{M} \omega_m = 1$ (Pawlikowski and Chorowska, 2020).

Alternatively, stacker weights ω can be learned from data by minimizing the loss function, $\omega = \arg\min_{\omega' \in \mathbb{R}^M} \mathcal{L}\left(g(\hat{y}_1, \ldots, \hat{y}_M), y_{\text{out}}\right)$, often subject to $\omega_m \geq 0$ and $\sum_m \omega_m = 1$. If the loss is differentiable, weights can be learned via a **numerical optimization** method such as gradient descent. Another popular approach is the greedy **ensemble selection** algorithm (Caruana et al., 2004), which adds models to the ensemble with replacement. Ensemble selection is used by default in several AutoML forecasting frameworks (Deng et al., 2022; Shchur et al., 2023; Zöller et al., 2024).

General combinations via linear regression. Instead of assigning a single weight per model, weights can vary with item i, time step h, quantile q, or any combination thereof. Reintroducing the item index i, we write:

$$[g_i(\hat{y}_{i,1}, \ldots, \hat{y}_{i,M})]_{h,q} = \sum_{m=1}^{M} \omega_{i,h,q,m} \cdot \left[\hat{y}_{i,m}\right]_{h,q} \tag{2}$$

with weights $\omega \in \mathbb{R}^{N \times H \times Q \times M}$. These weights can be tied across horizon steps, e.g., $\omega_{i,h,q,m} = \omega_{i,q,m}$, or even shared across items. This defines a broad family of linear combinations, as explored by Hasson et al. (2023) through different parameter-tying and regularization schemes. Our evaluation in Sec. 6 includes several of these linear variants.

Nonlinear regression. Finally, g can be a nonlinear function, often implemented using tree-based models such as gradient boosting—e.g., as in `darts` (Herzen et al., 2022), `sktime` (Löning et al., 2019), or Gastinger et al. (2021). As with linear models, these can vary in how they tie parameters across items, quantiles, or forecast horizons. Input normalization (e.g., scaling $\hat{y}_m$) is often beneficial before applying nonlinear models. In Sec. 6, we evaluate several nonlinear tabular models.

3.2 Training stacker models with time series cross-validation

One important detail that we have not yet addressed is how to obtain the training data for the stacker models. It is conventional wisdom in ensemble training that in order to prevent overfitting, stacker models g should be trained on *out-of-fold* data—i.e., data that was held out during the training of the base models $\{f_m\}$ (Zhou, 2012). Holding out data, in turn, requires special treatment in the case of time series (R. Hyndman and Athanasopoulos, 2021, Section 5.10). In line with this, we use a time series K-fold cross validation scheme which works as follows. For each fold $k = 1, \ldots, K$, we first remove the last $j := (K - k + 1)$ windows of size H from each time series, obtaining the training set $y_{1:T-jH}$ used to train the base models $\{f_m\}$. Next, we make H-step-ahead predictions with all trained base models, obtaining M quantile forecasts $\{\hat{y}_m^k\}$, each covering the validation window $y^k = y_{T-jH:T-(j-1)H}$. After repeating this process K times, we collect the training set for the stacker model g consisting of the inputs (the out-of-fold predictions $\hat{y}_m^k$ for each model m and fold k), and the targets (the out-of-fold data y^k for each fold k). See Fig. 1 for a visual depiction.

4 Multi-layer stacking for time series forecasting

4.1 Multi-layer stacking

As we will see in Sec. 6, the performance of different stacker models can vary greatly across datasets. To address this issue, we propose a *multi-layer stacking* approach that aggregates the outputs of multiple stackers with yet another stacker model (Fig. 2). This approach removes the need to decide on a single stacker model g in advance—we can instead learn the optimal combination of multiple stacker models $\{g_c\}_{c=1}^C$, similar to how we considered a set of base forecasting models $\{f_m\}_{m=1}^M$. Finally, we can combine the outputs of $\{g_c\}_{c=1}^C$ using a new stacker model $s\colon (\mathbb{R}^{H \times Q})^C \to \mathbb{R}^{H \times Q}$. In the following, we refer to the base models $\{f_m\}_{m=1}^M$ as the first layer of the multi-layer stack ensemble, the stacker models $\{g_c\}_{c=1}^C$ as the second layer, and the aggregator model s as the third layer; or in short as L1, L2, and L3 models (*cf.* Fig. 2).
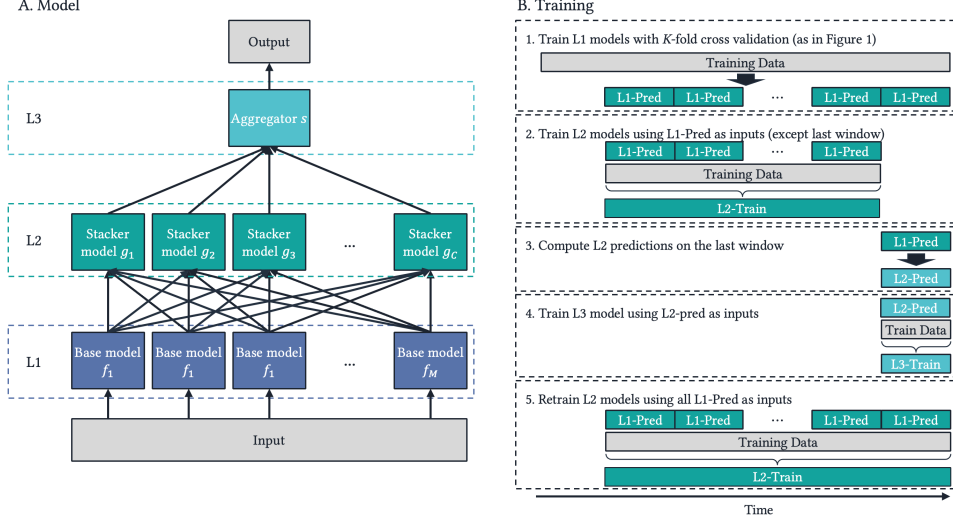
Figure 2: Architecture and training procedure of a multi-layer stacker model.

Concretely, to generate quantile predictions for an item, its past data $y_{i,1:T}$ is used to produce forecasts $\hat{y}_{i,1}, \ldots, \hat{y}_{i,M}$ from each L1 model. These are then aggregated by a composition of L2 models and a single L3 model s. The resulting *multi-layer stack ensemble* defines the mapping:

$$y_{i,1:T} \mapsto s(g_1(\hat{y}_{i,1}, \ldots, \hat{y}_{i,M}), \ldots, g_C(\hat{y}_{i,1}, \ldots, \hat{y}_{i,M})).$$

Note that all forecast combination and stacker models g discussed in the previous sections can also be used as the aggregator model s. For example, the aggregator s can perform model selection among the L2 stacker models, or it could combine the L2 models into a weighted ensemble.

4.2 Training of multi-layer stacker models with two-level time series cross-validation

Similar to the individual stacker models g, we need to make sure that the L3 aggregator model s is trained on out-of-fold predictions of the L2 models to avoid overfitting. To achieve this, we adjust the procedure described in Sec. 3.2. Like before, we train the L1 base models using K-fold time series cross validation. Next, we only train the L2 stacker models $\{g_c\}$ on the first $K-1$ validation windows. We then make predictions with the L2 models for the Kth validation window, and use these predictions together with the ground truth values to fit the L3 aggregator model s. Finally, we re-train the L2 models using all K validation windows to ensure that these models have access to the most recent data. The entire process is shown in Fig. 2.

5 Related work

Time series forecasting. Time series modeling spans classical statistical methods (Box et al., 1970; R. Hyndman, Koehler, et al., 2008), deep learning models (Benidis et al., 2022; Challu et al., 2023; Nie et al., 2022; Salinas, Flunkert, et al., 2020), and more recently, pretrained models (Ansari et al., 2024; Das et al., 2024; Woo et al., 2024). Despite substantial progress, no single model consistently outperforms others across all datasets and problem settings (Aksu et al., 2024; Godahewa et al., 2021). This motivates AutoML systems that automatically train, tune, and combine models to achieve the best performance for a given task (Ali, 2020; Deng et al., 2022; Shchur et al., 2023; Zöller et al., 2024). Our stacking framework is compatible with arbitrary forecasting models, making it complementary to ongoing advances in model development.

Forecast combinations. The idea of combining forecasts dates back to Bates and Granger (1969) and has inspired many methods (X. Wang et al., 2023), often featured in winning solutions of prediction competitions (Bojer and Meldgaard, 2021; Makridakis et al., 2018, 2022). Early studies found that simple averages often outperform more complex methods—a phenomenon known as the "forecast combination puzzle" (Smith and Wallis, 2009; Stock and Watson, 2004). However, these results were based on small datasets, statistical models, and point forecasts (Gastinger et al., 2021). Later work showed that more sophisticated methods outperform simple averages when applied to modern ML models, larger datasets, and probabilistic forecasting (Hasson et al., 2023). Our results in Sec. 6 advance this discussion and show that stacking improves over simple aggregation strategies for both point and probabilistic forecasting.

Stacking. Stacking, or stacked generalization, was introduced by Wolpert (1992) and Breiman (1996), where a single stacker model combines predictions from base models to produce the final output (Van der Laan et al., 2007). Variants of stacking have been studied extensively, including for quantile regression (Fakoor et al., 2023). In AutoML systems, a common approach is ensemble selection (Caruana et al., 2004), popularized by Auto-Sklearn (Feurer et al., 2015). Multi-layer stack ensembles first appeared in competition-winning solutions (Itericz and Semenov, 2016; Koren, 2009), and gained broader adoption with AutoGluon-Tabular (Erickson et al., 2020), which introduced a robust bagging-based implementation. However, existing work on multi-layer stacking is limited to tabular tasks. Our paper extends this framework to point and probabilistic time series forecasting.

6 Experiments

The main goal of our experimental analysis is to determine which forecast combination methods result in the best prediction accuracy (Sec. 6.1). In subsequent experiments, we aim to get better understanding of these results. For this purpose, we investigate the effects of various design choices such as selection of L1 models or the amount of validation data (Sec. 6.2–6.5).

6.1 Which combination methods produce the most accurate forecasts?

In our main experiment, we perform a large-scale benchmark comparison of 33 forecast combination methods, including the single-layer (Sec. 3) and multi-layer stacking (Sec. 4).

Datasets. We use 50 univariate datasets from Ansari et al. (2024) and Woo et al. (2024), covering diverse domains and frequencies, totaling 90K time series with 110M observations (see Tab. 5). To ensure that enough data is available for training base and stacker models, we only keep time series with at least $8 \times H$ observations (where H is the forecast horizon) in each dataset. As this filtering changes the data compared to the original publications, we re-evaluate all models ourselves.

Models. We consider 11 base forecasters (**L1 models**) covering all popular model categories: statistical models (*SeasonalNaive, AutoETS, Theta*), deep learning models (*DeepAR, PatchTST, TFT, TiDE, DLinear*), gradient-boosted trees (*DirectTabular, RecursiveTabular*), and one pretrained model (*Chronos-Bolt*). We use model implementations from AutoGluon–TimeSeries (Shchur et al., 2023), trained with default hyperparameters until convergence, without hyperparameter tuning for simplicity. These 11 L1 models are used as input to all the subsequent combination methods.

In addition, we train 31 combination methods (**L2 models**) from Sec. 3. These include 2 *simple averages, model selection*, 3 *performance-weighted averages*, 3 *greedy ensembles*, 18 *linear models*, and 4 *nonlinear models*. A detailed breakdown for each model category is available in Sec. B.

Finally, we consider 2 multi-layer stacking approaches (**L3 models**) from Sec. 4: *stacker model selection* that selects the best L2 model based on the performance on the last validation window, and *multi-layer stacking* that fits a greedy ensemble on top of the L2 models. To keep the training time reasonable, we limit the second level of the ensemble to 14 a priori chosen L2 models.

Table 1: Aggregated probabilistic forecasting performance of the representative combination methods based on SQL. Best result in **bold**, second best underlined. Individual results for all methods and datasets are available in Tab. 8

Method	(↑) Elo	(↑) Champion	(↓) Average rank	(↓) Average relative error	(↓) Median marginal training time
Median	1000	3	5.92	1.000	**1s**
Model selection	1049	<u>9</u>	5.43	1.001	<u>2s</u>
Performance-based average	1130	3	4.60	0.963	<u>2s</u>
Greedy ensemble selection	1191	4	3.92	0.952	11s
Linear model	<u>1220</u>	3	<u>3.62</u>	<u>0.947</u>	11s
Nonlinear model	1046	4	5.43	1.011	27s
Stacker model selection	1157	1	4.27	0.973	484s
Multi-layer stacking	**1306**	**20**	**2.81**	**0.945**	721s

Table 2: Aggregated probabilistic forecasting performance of the representative combination methods based on MASE. Best result in **bold**, second best underlined. Individual results for all methods and datasets are available in Tab. 9.

Method	(↑) Elo	(↑) Champion	(↓) Average rank	(↓) Average relative error	(↓) Median marginal training time
Median	1000	7	5.54	1.000	**1s**
Model selection	997	<u>8</u>	5.58	1.026	**1s**
Performance-based average	1059	2	4.98	0.989	**1s**
Greedy ensemble selection	1161	4	3.84	0.975	<u>2s</u>
Linear model	<u>1168</u>	3	<u>3.81</u>	0.975	5s
Nonlinear model	1004	5	5.48	1.033	6s
Stacker model selection	1158	3	3.87	<u>0.967</u>	295s
Multi-layer stacking	**1256**	**17**	**2.90**	**0.954**	443s

Metrics. For each dataset, we evaluate two tasks: point and probabilistic forecasting. Point accuracy is measured using *mean absolute scaled error* (MASE) on the median forecast, while probabilistic accuracy is assessed via *scaled quantile loss* (SQL) at quantile levels $\mathcal{Q} = \{0.1, 0.2, ..., 0.9\}$. To aggregate results across datasets, we report the *Elo rating*, number of wins (*champion*), *average rank*, and *average relative error* (geometric mean of errors normalized against the baseline). We use "Simple average (median)" as the baseline for Elo and relative error. Further details are provided in Sec. C.

Setup. We reserve the last H observations of each time series as the test set. L1 models are trained using K=5-fold cross-validation, with the model from the last fold used for test prediction. Individual L2 models are trained on all 5 validation windows. For multi-layer stacking, the L3 model is fitted on the final window. After training, all models generate forecasts on the test set, and accuracy is evaluated using MASE or SQL, depending on the task.

Results. Complete results for 2 task types × 50 datasets × 44 models are reported in Tables 8–9 (appendix). To summarize these, we group forecast combination methods into 8 categories: simple average, model selection, performance-based average, greedy ensemble selection, linear model, nonlinear model, stacker model selection, and multi-layer stacking. We select the best-performing method (by average rank) from each category as its *representative method* and report aggregate metrics for these in Tables 1–2. Note that such aggregation does not give an unfair advantage to the multi-layer stacking approaches since there is only one method in each category. Moreover, multi-layer stacking has the best scores both in the full results (Tables 8–9) and in the aggregated results (Tables 1–2). Based on these results, we draw the following conclusions:
1. **Combination methods outperform individual forecasting models**. Combining multiple forecasts typically improves the accuracy compared to a single model, even if model selection is

performed. We see major accuracy gains, with up to 200 Elo points difference ($\approx 75\%$ win rate) and up to 5% error reduction. This result reaffirms the previous findings on the importance of ensembling in forecasting (X. Wang et al., 2023).

2. **Stacking outperforms simple aggregation techniques**. Contrary to the "forecast combination puzzle," which suggests that complex methods rarely outperform simple ones, our results clearly demonstrate that learned aggregation methods—like ensemble selection and linear models—yield substantially better accuracy than simple or performance-based averages.

3. **Multi-layer stacking outperforms individual combination methods**. Across both point and probabilistic tasks, multi-layer stacking achieves the highest accuracy. This result is consistent with our earlier observations: since no single L2 model performs best across all datasets, combining them in an L3 ensemble leads to improved overall accuracy. Notably, multi-layer stacking also outperforms model selection over L2 models, highlighting that combining multiple stackers is more effective than selecting the best one in isolation.

6.2 When and why does multi-layer stacking outperform other approaches?

L3 model weights. First, we investigate the weights assigned by the L3 ensemble selection algorithm to the underlying L2 models. We show the average weights in Fig. 3 and provide the per-dataset breakdown in Fig. 9–10 in the appendix. All constituent L2 stacker models receive non-zero weight, underscoring the value of maintaining a diverse portfolio. Notably, the nonlinear stacker LightGBM performs poorly on average when evaluated in isolation (Tab. 1), yet it is frequently selected by the L3 ensemble. This indicates that while LightGBM may underperform overall, it excels on specific datasets—which can be capitalized on by the adaptive multi-layer stacking framework.

Figure 3: Weights assigned by the L3 ensemble selection algorithm to the L2 models (average over 50 tasks).

Normalized performance. To understand where multi-layer stacking succeeds or fails, we compare each representative model's score to the dataset-specific "champion" (Tab. 10). In 39 out of 50 datasets, most L2 models perform within 20% of the champion, and in these cases, multi-layer stacking typically ranks first or second in accuracy. In the remaining 11 datasets, where L2 model performance varies widely, multi-layer stacking tends to underperform. We hypothesize that allocating more validation data to train the L3 aggregator could mitigate this issue.

6.3 How much validation data do different combination methods require?

Previous experiments used 5 validation windows, requiring each base model to be trained 5 times. To reduce computational cost, we evaluated ensemble performance with fewer validation folds ($K = 1$ to 5), using Median aggregation as a baseline since its performance is fold-independent.

Results show that even with just $K = 2$, multi-layer stacking outperforms all other methods, followed by the linear model. The nonlinear model benefits most from additional folds, while model selection performs best at $K = 1$, suggesting that older data is less indicative of the test set performance.

Figure 4: Influence of the number of validation windows on the model performance.

Table 3: Ablation: Aggregated probabilistic forecasting performance of representative combination methods when using only 6 base models. Best result in **bold**, second best <u>underlined</u>. Multi-layer stacking remains the top performer even with a reduced set of base models.

Method	(↑) Elo	(↑) Champion	(↓) Average rank	(↓) Average relative error	(↓) Median marginal training time
Median	1000	1	5.38	1.000	**1s**
Model selection	899	2	6.39	1.058	**1s**
Performance-based average	1095	7	4.46	0.994	<u>2s</u>
Greedy ensemble selection	1110	6	4.23	0.990	7s
Linear model	<u>1145</u>	2	<u>3.93</u>	<u>0.978</u>	8s
Nonlinear model	1047	<u>8</u>	4.93	1.002	21s
Stacker model selection	1128	2	4.07	0.982	418s
Multi-layer stacking	**1280**	**22**	**2.61**	**0.947**	578s

Table 4: Ablation: Aggregated probabilistic forecasting performance of multi-layer stacking with and without L2 model retraining. Scores are computed with respect to the methods in Tab. 1. Skipping retraining reduces fit time, but decreases forecast accuracy.

Method	(↑) Elo	(↑) Champion	(↓) Average rank	(↓) Average relative error	(↓) Median marginal training time
Multi-layer stacking (L2 retraining)	**1306**	**20**	**2.81**	**0.945**	720s
Multi-layer stacking (no L2 retraining)	1253	13	3.24	0.950	**483s**

6.4 What is the effect of the L1 model choice?

Our experiments use a fixed set of 11 L1 forecasting models, but a robust ensembling method should perform well regardless of the base models used. To test this, we removed the 5 best-performing L1 models (*Chronos*, *PatchTST*, *TFT*, *DirectTabular*, *RecursiveTabular*) and retained only 6: *SeasonalNaive*, *AutoETS*, *Theta*, *DLinear*, *DeepAR*, and *TiDE*. The rest of the setup follows Sec. 6.1.

Aggregate results for 8 representative ensemble methods are shown in Tab. 3. The ranking remains unchanged: multi-layer stacking still performs best, showing robustness to L1 model choice. Model selection drops significantly due to the absence of Chronos (the strongest individual model), highlighting the importance of forecast combination when strong base models are unavailable.

6.5 Is it necessary to retrain the L2 models?

Our training procedure for the multi-layer stack ensemble (Sec. 4.2) includes a final step where L2 models are retrained on all validation windows. To assess whether this step is necessary, we perform an ablation comparing performance with and without this retraining. As shown in Tab. 4, skipping retraining leads to 1.5x faster training at the cost of lower forecast accuracy. This highlights the importance of retraining L2 models on the full validation data to ensure optimal performance.

7 Discussion

Limitations & future work. While our study demonstrates the effectiveness of learned forecast combination methods, several limitations remain. Most stacker models require predictions from all base models, which results in slower inference compared to sparse methods like ensemble selection. Future work could explore pruning strategy to mitigate this issue (Tsoumakas et al., 2009). Similarly, multi-layer stacking incurs substantial training costs, as we currently train 14 L2 models. A more principled approach—such as offline portfolio optimization, akin to Salinas and Erickson (2024)—could reduce computational overhead while maintaining or even improving accuracy. In addition, this framework can be extended with meta-learning approaches such as FFORMA (Montero-Manso et al., 2020), which aim to select or weight models based on dataset-level

features. Finally, the search for better aggregation models to use at the L2 and L3 levels remains an open and important challenge.

Summary. Our large-scale empirical study shows that learned ensembling methods—particularly stacking—consistently outperform both individual forecasting models and simple aggregation techniques. These findings challenge the "forecast combination puzzle" and demonstrate that flexible, data-driven ensemble strategies can significantly improve predictive accuracy. Multi-layer stacking extends this idea by combining multiple stackers, yielding robust performance across a wide range of datasets. While this approach provides consistent improvements overall, our ablations show it is especially valuable when strong individual forecasting models are unavailable. As forecasting tools and AutoML systems continue to evolve, we see strong potential in integrating advanced ensemble methods that adapt to data characteristics, scale efficiently, and generalize well across diverse forecasting scenarios.

8 Broader Impact Statement

After careful reflection, the authors have determined that this work presents no notable negative impacts to society or the environment.

References

Adhikari, R. and R. Agrawal (2012). "Combining multiple time series models through a robust weighted mechanism". In: *2012 1st International Conference on Recent Advances in Information Technology (RAIT)*. IEEE, pp. 455–460.

Aksu, T. et al. (2024). "GIFT-Eval: A Benchmark For General Time Series Forecasting Model Evaluation". In: *arxiv preprint arxiv:2410.10393*.

Ali, M. (Apr. 2020). *PyCaret: An open source, low-code machine learning library in Python*. PyCaret version 1.0.0. URL: https://www.pycaret.org.

Ansari, A. F. et al. (2024). "Chronos: Learning the language of time series". In: *arXiv preprint arXiv:2403.07815*.

Assimakopoulos, V. and K. Nikolopoulos (2000). "The theta model: a decomposition approach to forecasting". In: *International journal of forecasting* 16.4, pp. 521–530.

Bai, Y. et al. (2022). *Training a Helpful and Harmless Assistant with Reinforcement Learning from Human Feedback*. arXiv: 2204.05862 [cs.CL]. URL: https://arxiv.org/abs/2204.05862.

Bates, J. M. and C. W. Granger (1969). "The combination of forecasts". In: *Journal of the operational research society* 20.4, pp. 451–468.

Benidis, K. et al. (2022). "Deep learning for time series forecasting: Tutorial and literature survey". In: *ACM Computing Surveys* 55.6, pp. 1–36.

Bojer, C. S. and J. P. Meldgaard (2021). "Kaggle forecasting competitions: An overlooked learning opportunity". In: *International Journal of Forecasting* 37.2, pp. 587–603.

Boubdir, M. et al. (2024). "Elo Uncovered: Robustness and Best Practices in Language Model Evaluation". In: *The Thirty-eighth Annual Conference on Neural Information Processing Systems*. URL: https://openreview.net/forum?id=Pc9LLjTL5f.

Box, G. E. et al. (1970). *Time series analysis: forecasting and control*. John Wiley & Sons.

Breiman, L. (1996). "Stacked regressions". In: *Machine learning* 24, pp. 49–64.

Caruana, R. et al. (2004). "Ensemble selection from libraries of models". In: *Proceedings of the twenty-first international conference on Machine learning*, p. 18.

Challu, C. et al. (2023). "Nhits: Neural hierarchical interpolation for time series forecasting". In: *Proceedings of the AAAI conference on artificial intelligence*. Vol. 37. 6, pp. 6989–6997.

Chiang, W.-L. et al. (2024). *Chatbot Arena: An Open Platform for Evaluating LLMs by Human Preference*. arXiv: 2403.04132 [cs.AI]. URL: https://arxiv.org/abs/2403.04132.

Clemen, R. T. (1989). "Combining forecasts: A review and annotated bibliography". In: *International journal of forecasting* 5.4, pp. 559–583.

Croston, J. D. (Sept. 1972). "Forecasting and Stock Control for Intermittent Demands". In: *Journal of the Operational Research Society* 23.3, pp. 289–303. URL: https://doi.org/10.1057/jors.1972.50.

Das, A. et al. (2023). "Long-term forecasting with TiDE: Time-series dense encoder". In: *arXiv preprint arXiv:2304.08424*.

Das, A. et al. (2024). "A decoder-only foundation model for time-series forecasting". In: *Proceedings of the 41st International Conference on Machine Learning*. Ed. by R. Salakhutdinov et al. Vol. 235. Proceedings of Machine Learning Research. PMLR, pp. 10148–10167. URL: https://proceedings.mlr.press/v235/das24c.html.

Deng, D. et al. (2022). "Efficient automated deep learning for time series forecasting". In: *Joint European Conference on Machine Learning and Knowledge Discovery in Databases*. Springer, pp. 664–680.

Ekambaram, V. et al. (2024). *Tiny Time Mixers (TTMs): Fast Pre-trained Models for Enhanced Zero/Few-Shot Forecasting of Multivariate Time Series*. arXiv: 2401.03955 [cs.LG]. URL: https://arxiv.org/abs/2401.03955.

Elo, A. E. (1967). "The proposed uscf rating system, its development, theory, and applications". In: *Chess life* 22.8, pp. 242–247.

Erickson, N. et al. (2020). "Autogluon-tabular: Robust and accurate automl for structured data". In: *arXiv preprint arXiv:2003.06505*.

Fakoor, R. et al. (2023). "Flexible model aggregation for quantile regression". In: *Journal of Machine Learning Research* 24.162, pp. 1–45.

Feurer, M. et al. (2015). "Efficient and robust automated machine learning". In: *Advances in neural information processing systems* 28.

Fiorucci, J. A. et al. (2016). "Models for optimising the theta method and their relationship to state space models". In: *International journal of forecasting* 32.4, pp. 1151–1161.

Garza, A. et al. (2022). *StatsForecast: Lightning fast forecasting with statistical and econometric models*. PyCon Salt Lake City, Utah, US 2022. URL: https://github.com/Nixtla/statsforecast.

Gastinger, J. et al. (2021). "A study on ensemble learning for time series forecasting and the need for meta-learning". In: *2021 International Joint Conference on Neural Networks (IJCNN)*. IEEE, pp. 1–8.

Gijsbers, P. et al. (2024). "Amlb: an automl benchmark". In: *Journal of Machine Learning Research* 25.101, pp. 1–65.

Godahewa, R. et al. (2021). "Monash Time Series Forecasting Archive". In: *Neural Information Processing Systems Track on Datasets and Benchmarks*.

Hasson, H. et al. (2023). "Theoretical guarantees of learning ensembling strategies with applications to time series forecasting". In: *International Conference on Machine Learning*. PMLR, pp. 12616–12632.

Herzen, J. et al. (2022). "Darts: User-Friendly Modern Machine Learning for Time Series". In: *Journal of Machine Learning Research* 23.124, pp. 1–6. URL: http://jmlr.org/papers/v23/21-1177.html.

Holzmüller, D., L. Grinsztajn, and I. Steinwart (2024). "Better by default: Strong pre-tuned mlps and boosted trees on tabular data". In: *Advances in Neural Information Processing Systems* 37, pp. 26577–26658.

Hong, T. et al. (2016). "Probabilistic energy forecasting: Global Energy Forecasting Competition 2014 and beyond". In: *International Journal of Forecasting* 32.3, pp. 896–913. URL: https://www.sciencedirect.com/science/article/pii/S0169207016000133.

Hyndman, R., A. B. Koehler, et al. (2008). *Forecasting with exponential smoothing: the state space approach*. Springer Science & Business Media.

Hyndman, R. J. and Y. Khandakar (2008). "Automatic time series forecasting: the forecast package for R". In: *Journal of statistical software* 27, pp. 1–22.

Hyndman, R. and G. Athanasopoulos (2021). *Forecasting: Principles and Practice*. English. 3rd. Australia: OTexts.

Itericz, G. and S. Semenov (2016). *1st place - winner solution. Otto Group Product Classification Challenge*. https://www.kaggle.com/c/otto-group-product-classification-challenge/discussion/14335. Accessed: 2025-03-21.

Januschowski, T. et al. (2020). "Criteria for classifying forecasting methods". In: *International Journal of Forecasting* 36.1, pp. 167–177.

Ke, G. et al. (2017). "Lightgbm: A highly efficient gradient boosting decision tree". In: *Advances in neural information processing systems* 30.

Kingma, D. P. (2014). "Adam: A method for stochastic optimization". In: *arXiv preprint arXiv:1412.6980*.

Koren, Y. (2009). "The bellkor solution to the netflix grand prize". In: *Netflix prize documentation* 81.2009, pp. 1–10.

Lim, B. et al. (2021). "Temporal fusion transformers for interpretable multi-horizon time series forecasting". In: *International Journal of Forecasting* 37.4, pp. 1748–1764.

Löning, M. et al. (2019). "sktime: A unified interface for machine learning with time series". In: *arXiv preprint arXiv:1909.07872*.

Makridakis, S., E. Spiliotis, and V. Assimakopoulos (2018). "The M4 Competition: Results, findings, conclusion and way forward". In: *International Journal of forecasting* 34.4, pp. 802–808.

Makridakis, S., E. Spiliotis, and V. Assimakopoulos (2022). "M5 accuracy competition: Results, findings, and conclusions". In: *International Journal of Forecasting* 38.4, pp. 1346–1364.

Montero-Manso, P. et al. (2020). "FFORMA: Feature-based forecast model averaging". In: *International Journal of Forecasting* 36.1, pp. 86–92.

Nie, Y. et al. (2022). "A time series is worth 64 words: Long-term forecasting with transformers". In: *arXiv preprint arXiv:2211.14730*.

Pawlikowski, M. and A. Chorowska (2020). "Weighted ensemble of statistical models". In: *International Journal of Forecasting* 36.1, pp. 93–97.

Perrone, M. P. and L. N. Cooper (1995). "When networks disagree: Ensemble methods for hybrid neural networks". In: *How we learn; How we remember: Toward an understanding of brain and neural systems: Selected papers of Leon N Cooper*. World Scientific, pp. 342–358.

Purucker, L. O. and J. Beel (2023). "CMA-ES for post hoc ensembling in AutoML: a great success and salvageable failure". In: *International Conference on Automated Machine Learning*. PMLR, pp. 1–1.

Salinas, D. and N. Erickson (2024). *TabRepo: A Large Scale Repository of Tabular Model Evaluations and its AutoML Applications*. arXiv: 2311.02971 [cs.LG]. URL: https://arxiv.org/abs/2311.02971.

Salinas, D., V. Flunkert, et al. (2020). "DeepAR: Probabilistic forecasting with autoregressive recurrent networks". In: *International journal of forecasting* 36.3, pp. 1181–1191.

Shchur, O. et al. (2023). "AutoGluon–TimeSeries: AutoML for probabilistic time series forecasting". In: *International Conference on Automated Machine Learning*. PMLR, pp. 9–1.

Smith, J. and K. F. Wallis (2009). "A simple explanation of the forecast combination puzzle". In: *Oxford bulletin of economics and statistics* 71.3, pp. 331–355.

Stock, J. H. and M. W. Watson (2004). "Combination forecasts of output growth in a seven-country data set". In: *Journal of forecasting* 23.6, pp. 405–430.

Tsang, T. K. et al. (2024). "An adaptive weight ensemble approach to forecast influenza activity in an irregular seasonality context". In: *Nature Communications* 15.1, p. 8625. URL: https://doi.org/10.1038/s41467-024-52504-1.

Tsoumakas, G., I. Partalas, and I. Vlahavas (2009). "An ensemble pruning primer". In: *Applications of supervised and unsupervised ensemble methods*, pp. 1–13.

Van der Laan, M. J., E. C. Polley, and A. E. Hubbard (2007). "Super learner". In: *Statistical applications in genetics and molecular biology* 6.1.

Wang, X. et al. (2023). "Forecast combinations: An over 50-year review". In: *International Journal of Forecasting* 39.4, pp. 1518–1547.

Wolpert, D. H. (1992). "Stacked generalization". In: *Neural networks* 5.2, pp. 241–259.

Woo, G. et al. (2024). *Unified Training of Universal Time Series Forecasting Transformers*. arXiv: 2402.02592 [cs.LG]. URL: https://arxiv.org/abs/2402.02592.

Zeng, A. et al. (2023). "Are transformers effective for time series forecasting?" In: *Proceedings of the AAAI conference on artificial intelligence*. Vol. 37. 9, pp. 11121–11128.

Zhou, Z.-H. (2012). *Ensemble methods: foundations and algorithms*. CRC press.

Zöller, M.-A., M. Lindauer, and M. F. Huber (2024). "Auto-sktime: Automated Time Series Forecasting". In: *International Conference on Learning and Intelligent Optimization*. Springer, pp. 456–471.

Submission Checklist

1. For all authors...

   (a) Do the main claims made in the abstract and introduction accurately reflect the paper's contributions and scope? [Yes]

   (b) Did you describe the limitations of your work? [Yes] See Sec. 7.

   (c) Did you discuss any potential negative societal impacts of your work? [N/A] After careful consideration, we could not identify any notable negative societal impacts of our work.

   (d) Did you read the ethics review guidelines and ensure that your paper conforms to them? `https://2022.automl.cc/ethics-accessibility/` [Yes]

2. If you ran experiments...

   (a) Did you use the same evaluation protocol for all methods being compared (e.g., same benchmarks, data (sub)sets, available resources)? [Yes] All evaluated models and methods were implemented in the same software framework, ensuring uniform evaluation.

   (b) Did you specify all the necessary details of your evaluation (e.g., data splits, pre-processing, search spaces, hyperparameter tuning)? [Yes] All relevant details are provided in Sec. 6 and the appendix.

   (c) Did you repeat your experiments (e.g., across multiple random seeds or splits) to account for the impact of randomness in your methods or data? [No] Due to a large number of evaluated models and datasets (44 models × 50 datasets × 2 task types), we evaluated all models using a single random seed.

   (d) Did you report the uncertainty of your results (e.g., the variance across random seeds or splits)? [Yes] In addition to the aggregated scores reported in Tables1–2, we visualize the distribution of the relative model errors in Sec. E.2.

   (e) Did you report the statistical significance of your results? [Yes] We include the critical differences diagrams in Sec. E.

   (f) Did you use tabular or surrogate benchmarks for in-depth evaluations? [No] No tabular or surrogate benchmarks are available for the tasks considered in this paper.

   (g) Did you compare performance over time and describe how you selected the maximum duration? [N/A] We provided a fixed training time budget to all methods, capping the runtime for individual L2 ensemble methods to 10 minutes.

   (h) Did you include the total amount of compute and the type of resources used (e.g., type of GPUs, internal cluster, or cloud provider)? [Yes] We include information about resources and training times in Sec. D

   (i) Did you run ablation studies to assess the impact of different components of your approach? [Yes] We include numerous ablation studies in Sec. 6.

3. With respect to the code used to obtain your results...

   (a) Did you include the code, data, and instructions needed to reproduce the main experimental results, including all requirements (e.g., `requirements.txt` with explicit versions), random seeds, an instructive `README` with installation, and execution commands (either in the supplemental material or as a URL)? [Yes] Included in the supplementary materials.

(b) Did you include a minimal example to replicate results on a small subset of the experiments or on toy data? [Yes]

(c) Did you ensure sufficient code quality and documentation so that someone else can execute and understand your code? [Yes]

(d) Did you include the raw results of running your experiments with the given code, data, and instructions? [Yes] Raw results are available in `results/` in the supplementary materials.

(e) Did you include the code, additional data, and instructions needed to generate the figures and tables in your paper based on the raw results? [Yes] Aggregation logic for the results is included in the supplementary materials.

4. If you used existing assets (e.g., code, data, models)…

(a) Did you cite the creators of used assets? [Yes]

(b) Did you discuss whether and how consent was obtained from people whose data you're using/curating if the license requires it? [N/A] We only used publicly available datasets with open source licenses that don't contain PII.

(c) Did you discuss whether the data you are using/curating contains personally identifiable information or offensive content? [N/A] The datasets used for evaluation contain no PII.

5. If you created/released new assets (e.g., code, data, models)…

(a) Did you mention the license of the new assets (e.g., as part of your code submission)? [N/A] Our experiments were conducted on publicly available datasets and we have not introduced new datasets.

(b) Did you include the new assets either in the supplemental material or as a URL (to, e.g., GitHub or Hugging Face)? [N/A] Our experiments were conducted on publicly available datasets and we have not introduced new datasets.

6. If you used crowdsourcing or conducted research with human subjects…

(a) Did you include the full text of instructions given to participants and screenshots, if applicable? [N/A] We did not use crowdsourcing.

(b) Did you describe any potential participant risks, with links to Institutional Review Board (IRB) approvals, if applicable? [N/A] We did not use crowdsourcing.

(c) Did you include the estimated hourly wage paid to participants and the total amount spent on participant compensation? [N/A] We did not use crowdsourcing.

7. If you included theoretical results…

(a) Did you state the full set of assumptions of all theoretical results? [N/A] No theoretical results were included.

(b) Did you include complete proofs of all theoretical results? [N/A] No theoretical results were included.

A Datasets

Table 5: Dataset statistics. Note that all datasets were filtered to only contain time series with at least $8 \times H$ observations to ensure that enough training & validation data is available.

Dataset	Freq.	Seasonality	Horizon (H)	Num. Series	Num. Obs.	Min. Length	Max. Length
BDG-2 Bear	h	24	48	91	1,482,312	8,760	17,544
BDG-2 Bull	h	24	48	41	719,304	17,544	17,544
BDG-2 Fox	h	24	48	135	2,324,568	8,760	17,544
BDG-2 Hog	h	24	48	24	421,056	17,544	17,544
BDG-2 Panther	h	24	48	105	919,800	8,760	8,760
BDG-2 Rat	h	24	48	280	4,728,288	8,760	17,544
Beijing Air Quality	h	24	48	132	4,628,448	35,064	35,064
Beijing Subway	30min	48	96	552	867,744	1,572	1,572
Borealis	h	24	48	15	83,269	3,528	7,447
CDC Fluview ILINet	W	1	8	375	319,515	211	1,359
Favorita Store Sales	D	7	16	1,782	3,008,016	1,688	1,688
Favorita Transactions	D	7	16	54	91,152	1,688	1,688
GEF12	h	24	48	11	433,554	39,414	39,414
GEF17	h	24	48	8	140,352	17,544	17,544
HZMetro	15min	96	96	160	380,320	2,377	2,377
Hierarchical Sales	D	7	14	118	212,164	1,798	1,798
IDEAL	h	24	48	217	1,255,253	393	16,167
KDD Cup 2022	10min	144	96	134	4,727,519	35,279	35,280
Los-Loop	5min	288	96	207	7,094,304	34,272	34,272
M-Dense	h	24	48	30	525,600	17,520	17,520
PEMS03	5min	288	96	358	9,382,464	26,208	26,208
PEMS08	5min	288	96	510	9,106,560	17,856	17,856
Project Tycho	W	1	8	1,258	1,377,707	102	3,854
SHMetro	15min	96	96	576	5,073,984	8,809	8,809
SMART	h	24	48	5	95,709	8,398	26,304
SZ-Taxi	15min	96	96	156	464,256	2,976	2,976
Subseasonal Precipitation	D	7	14	862	9,760,426	11,323	11,323
Australian Electricity	30min	48	96	5	1,155,264	230,736	232,272
ERCOT	h	24	24	8	1,238,976	154,872	154,872
ETT (15 Min.)	15min	96	96	14	975,520	69,680	69,680
ETT (Hourly)	h	24	24	14	243,880	17,420	17,420
Electricity (Hourly)	h	24	48	321	8,443,584	26,304	26,304
Electricity (Weekly)	W	1	8	321	50,076	156	156
FRED-MD	M	12	12	107	77,896	728	728
KDD Cup 2018	h	24	48	270	2,942,364	9,504	10,920
M4 (Daily)	D	1	14	4,218	10,022,845	112	9,933
M4 (Hourly)	h	24	48	414	373,372	748	1,008
M4 (Monthly)	M	12	18	32,436	9,773,903	144	2,812
M4 (Quarterly)	Q	4	8	19,049	2,167,663	64	874
M4 (Weekly)	W	1	13	294	365,534	260	2,610
NN5 (Daily)	D	7	56	111	87,801	791	791
NN5 (Weekly)	W	1	8	111	12,543	113	113
Pedestrian Counts	h	24	168	64	3,130,594	1,777	96,424
Solar (10 Min.)	10min	144	96	137	7,200,720	52,560	52,560
Taxi (30 Min.)	30min	48	96	2,428	3,589,798	1,469	1,488
Taxi (Hourly)	h	24	48	2,428	1,794,292	734	744
Traffic (Weekly)	W	1	8	862	89,648	104	104
Uber TLC (Daily)	D	7	14	262	47,422	181	181
Uber TLC (Hourly)	h	24	48	262	1,138,128	4,344	4,344
Wind Farms (Daily)	D	7	14	335	119,407	149	366

B Models

B.1 Base models

We considered the following base models, using their implementations provided by Auto-Gluon–TimeSeries (Shchur et al., 2023):

- `SeasonalNaive`: Simple model that sets the forecast equal to the last observed value from the same season (R. Hyndman and Athanasopoulos, 2021).

- `AutoETS`: Automatically tuned exponential smoothing with trend and seasonality (Garza et al., 2022; R. Hyndman, Koehler, et al., 2008; R. J. Hyndman and Khandakar, 2008)

- `DynamicOptimizedTheta`: A generalization of the Theta method by Assimakopoulos and Nikolopoulos (2000), that automatically selects and revises some of the model hyperparameters (Fiorucci et al., 2016).

- `DeepAR`: Autoregressive forecasting model based on a recurrent neural network (Salinas, Flunkert, et al., 2020).

- `PatchTST`: Transformer-based forecaster that segments each time series into patches (Nie et al., 2022).

- `TemporalFusionTransformer`: Combines LSTM with a transformer layer to predict the quantiles of all future target values (Lim et al., 2021).

- `DirectTabular`: Predict the future time series values by transforming the task into a tabular prediction task and then applying the LightGBM quantile regressor (Ke et al., 2017).

- `RecursiveTabular`: Predict the future time series values one-by-one by transforming the task into a tabular prediction task and then applying the LightGBM regressor (Ke et al., 2017). In contrast to `DirectTabular`, the forecast is computed step-by-step by repeatedly applying the tabular method.

- `TiDEModel`: Time series dense encoder model (Das et al., 2023).

- `Chronos`: Pretrained time series forecasting model (Ansari et al., 2024). We use the `bolt_base` configuration.

For simplicity, we do not investigate the effect of hyperparameter tuning on the models and keep all the hyperparameters to their default values. We do not expect this to affect our main conclusions, given the stability of our results with respect to the L1 model choice (Sec. 6.4).

For each dataset, we generate the base model predictions by fitting all models with 5 validation windows, refitting each model from scratch for each window. We set the maximum time limit of 30 minutes per window per model to avoid extremely long runtimes, though the vast majority of models complete training within 5 minutes per window. We save the base model predictions for all validation windows and the test window to disk. This enables us to train the stacker models without needing to re-fit the base models.

B.2 Individual forecast combination methods

- **Simple averages**: We consider two simple averages, *mean* and *median*, which compute the mean or median of all base model predictions, respectively.

- **Model selection**: Model selection computes the index of the best-performing model according to the validation data during training, and then uses this individual model for forecasting.

- **Performance-weighted averages**: These models are weighted averages of the form $g(\hat{y}_1, \ldots, \hat{y}_M) = \sum_{m=1}^{M} \omega_m \hat{y}_m$, where the weights ω_m are computed directly from the (normalized) validation scores of the individual models, as $\omega_m \propto h(L_m)$, with $L_m \propto \mathcal{L}(\hat{y}_m, y_{\text{out}})$ such that $\sum_{m=1}^{M} L_m = 1$, and $\sum_{m=1}^{M} \omega_m = 1$. We consider three options for the function h:
  - Inv.: $h_{\text{inv}}(L) = 1/L$,
  - Sqr.: $h_{\text{sqr}}(L) = 1/L^2$,

– Exp.: $h_{\exp}(L) = \exp(1/L)$

Refer to Pawlikowski and Chorowska (2020) for more details.

- **Greedy ensembles**: This is the *ensemble selection* method by (Caruana et al., 2004), which optimizes the validation loss by greedily adding models to an equally-weighted ensemble *with replacement.* As an equally-weighted ensemble with replacement is equivalent to a weighted average with fractional weights, the greedy ensemble can also be interpreted as a weighted average $g(\hat{y}_1, \ldots, \hat{y}_M) = \sum_{m=1}^{M} \omega_m \hat{y}_m$ which is optimized via coordinate-wise ascent: Starting with zero weights $\omega^{(0)} = 0 \in \mathbb{R}^M$, the algorithm iterates for $j = 1, \ldots, S$, and selects at each step the model that would minimize the resulting loss, i.e.

$$m^{(j)} = \underset{m=1,\ldots M}{\arg\min} \, \mathcal{L}\left(g_{\hat{\omega}_m^{(j)}}(\hat{y}_1, \ldots, \hat{y}_M), y_{\text{ensemble}}\right), \tag{3}$$

where $\hat{\omega}_m^{(j)} \coloneqq ((j-1)\omega^{(j-1)} + e_m)/j$ are the weights that the greedy ensemble would have if it were to select model m, with $e_m \in \mathbb{R}^M$ being a canonical basis vector, i.e. $[e_m]_k = \mathbb{1}_{k=m}$. It then adds the model to the ensemble and sets $\omega^{(j)} = \hat{\omega}_{m^{(j)}}^{(j)}$. This approach is currently also used as the default in multiple AutoML forecasting frameworks (Deng et al., 2022; Shchur et al., 2023; Zöller et al., 2024). In our evaluation, we consider three versions of this model, with the number of iterations S set to 10, 100, and 1000, respectively.

- **Linear models**: For linear models, we consider a class of models g_i of the form

$$[g_i(\hat{y}_{i,1}, \ldots, \hat{y}_{i,M})]_{h,q} = \sum_{m=1}^{M} \omega_{i,h,q,m} \cdot [\hat{y}_{i,m}]_{h,q} \tag{4}$$

with weights $\omega \in \mathbb{R}^{N \times H \times Q \times M}$. More concretely, we consider a number of variations of this model which satisfy different constraints or have different degrees of weight-tying:

- *Positivity and simplex constraints:* It is often desirable to enforce positivity and/or simplex constraints, that is, $\omega_{i,h,q,m} > 0$ and $\sum_m \omega_{i,h,q,m} = 1$, respectively (X. Wang et al., 2023). In our model, we enforce constraints by parameterizing the weights through an unconstrained weight $\tilde{\omega}_{i,h,q,m}$ and then passing it through an activation function. We consider two versions:
  * `softmax`: This version enforces both positivity and simplex constraints by passing the unconstrained weights through a softmax activation function:

$$\omega_{i,h,q,m} = \frac{\exp(\tilde{\omega}_{i,h,q,m})}{\sum_{m'} \exp(\tilde{\omega}_{i,h,q,m'})}. \tag{5}$$

  * `positive`: This version enforces positivity, but not the simplex constraint, by passing the unconstrained weights through a quadratic function:

$$\omega_{i,h,q,m} = (\tilde{\omega}_{i,h,q,m})^2. \tag{6}$$

- *Weight-tying:* Instead of considering different weights for each model, item, time step, and quantile, we can lower the complexity by tying weights across some of these dimensions (Hasson et al., 2023). For example, tying weights across items would imply $\omega_{i,h,q,m} = \omega_{i',h,q,m}$ for all items i, i'. This yields a number of different linear models, denoted by the dimension along which the weights are *not* tied:
  * `m`: One weight per model; weights are tied across items, prediction times, and quantiles. This case is equivalent to weighted average models.
  * `mi`: One weight per model and per item; weights are tied across prediction times and quantiles.

* `mt`: One weight per model and per prediction time; weights are tied across items and quantiles.
* `mq`: One weight per model and per quantile; weights are tied across items and prediction times.
* `mit`: One weight per model, item, and prediction time; weights are tied across quantiles.
* `miq`: One weight per model, item, and quantile; weights are tied across prediction times.
* `mtq`: One weight per model, prediction time, and quantile; weights are tied across items.
* `mitq`: One weight per model, item, prediction time, and quantile; no weight-tying.

- *Across-quantile weights:* Instead of treating all quantiles independently, we can also compute quantile predictions *across quantiles* and rely on the base model estimates of all quantiles, as proposed by Fakoor et al. (2023); In its most flexible form, this corresponds to a model of the form

$$[g_i(\hat{y}_{i,1}, \ldots, \hat{y}_{i,M})]_{h,q} = \sum_{m=1}^{M} \sum_{q' \in \mathcal{Q}} \omega_{i,h,q,q',m} \cdot [\hat{y}_{i,m}]_{h,q'}, \tag{7}$$

with weights $\omega \in \mathbb{R}^{N \times H \times Q \times Q \times M}$. As before, we consider `softmax` and `positive` versions as well as different weight-tying options:
* `mqq`: One weight per model and across quantiles; weights are tied across items and prediction times.
* `miqq`: One weight per model, items, and across quantiles; weights are tied across prediction times.
* `mtqq`: One weight per model, prediction time, and across quantiles; weights are tied across items.

Training. We train linear stacker models by optimizing the resulting loss:

$$\omega = \underset{\omega' \in \mathbb{R}^{N \times H \times Q \times M}}{\arg\min} \sum_{i} \mathcal{L}\left(g_i(\hat{y}_{i,1}, \ldots, \hat{y}_{i,M}), y_{i,T+1:T+H}\right), \tag{8}$$

where g_i is a linear model that uses the to-be-optimized weights ω'. This can be approached with any numerical optimizer. In our implementation, we rely on the Adam optimizer (Kingma, 2014) and we use a custom learning rate schedule to reduce the learning rate whenever a plateau is hit and the loss starts oscillating. We also limit the training time to 10 minutes.

• **Nonlinear models**: To apply nonlinear models, we re-formulate the stacking problem as a tabular regression problem. Each item × timestep combination corresponds to a row in the training data. The quantile forecasts of the L1 models are used as features, and the ground truth time series value is the training target. We can then train any tabular regression model on this data. In our experiments we consider two methods:

- `RealMLP` is a deep-learning approach for tabular problems which improves on standard MLPs through a number of tricks and better, meta-learned default parameters (Holzmüller et al., 2024). We use the multi-quantile loss evaluated at levels $\mathcal{Q}$ as the training objective.
- `LightGBM` is a highly efficient implementation of gradient-boosted decision trees (Ke et al., 2017). We train a separate LightGBM regressor to predict each of the quantile levels $\mathcal{Q}$.

Additionally, we consider *scaled* versions of both methods in which we normalize the model predictions before feeding them into the tabular stacker model and then un-normalize the outputs again. This corresponds to a modified model g' of the form

$$g_i'(\hat{y}_{i,1}, \ldots, \hat{y}_{i,M}) = \frac{g_i(\alpha \hat{y}_{i,1} + \beta, \ldots, \alpha \hat{y}_{i,M} + \beta) - \beta}{\alpha}, \tag{9}$$

where α, β are computed from the empirical mean and standard deviation of the predictions in order to standardize the inputs, and g is a tabular stacker model as introduced above.

From each category we select one "representative model" for our condensed experiment summary in the paper, as described in Tab. 6 and Tab. 7.

Table 6: Representative models per category for the quantile forecast experiment.

Model category	Model
Median	Median
Model Selection	Model selection
Performance-based average	Performance-based average (exp)
Greedy ensemble selection	Greedy (S=100)
Linear model	Linear(mq, softmax)
Nonlinear model	LightGBM (scaled)

Table 7: Representative models per category for the point forecast experiment.

Model category	Model
Median	Median
Model Selection	Model selection
Performance-based average	Performance-based average (exp)
Greedy ensemble selection	Greedy (S=100)
Linear model	Linear(m, softmax)
Nonlinear model	LightGBM (scaled)

B.3 Multi-layer stacker models

For multilayer stacking, we use the following 14 stacker models as L2 models:

- Median
- Greedy (S=100)
- Linear (mi, softmax)
- Linear (mt, softmax)
- Linear (mq, softmax)
- Linear (mit, positive)
- Linear (mtq, positive)
- Linear (miq, positive)
- Linear (mqq, positive)
- Linear (miqq, positive)
- Linear (mtqq, positive)
- LightGBM
- LightGBM (scaled)
- RealMLP (scaled)

We have arbitrarily chosen these 14 L2 models since they provide a good coverage of different model families. To ensure a fair comparison, we have fixed this selection before running any experiments, and did not adjust the selection to optimize the benchmark performance.

We then consider multi-layer stackers resulting from two different L3 models:

- **Stacker model selection** uses *model selection* as its L3 model and thus relies on a single L2 model during test time.

- **Multi-layer stacking** uses *Greedy (S=100)* as its L3 model to compute a weighted average of the L2 model predictions.

C Evaluation metrics

C.1 Loss functions

Loss functions are used both to train and evaluate time series models. In Sec. 2.1, we introduced them as functions that take in the (quantile) prediction of a particular item as well as the ground-truth values and return a positive scalar value, that is

$$\mathcal{L} : \mathbb{R}^{H \times Q} \times \mathbb{R}^H \to \mathbb{R}_{\geq 0}, \tag{10}$$

with lower values indicating a more accurate forecast. We consider the following two losses to evaluate quantile forecasts and point forecasts, respectively:

- **Scaled quantile loss (SQL)**:

$$\mathrm{SQL}(\hat{y}_i, y_i) \coloneqq \frac{1}{H} \frac{1}{Q} \frac{1}{a_i} \sum_{h=1}^{H} \sum_{q \in \mathcal{Q}} \rho_q\left(\hat{y}_{i,T+h}^q, y_{i,T+h}\right), \tag{11}$$

where $\rho_q : \mathbb{R} \times \mathbb{R} \to \mathbb{R}$ is the *quantile loss at level q*, defined as

$$\rho_q\left(\hat{y}_{i,h}^q, y_{i,h}\right) \coloneqq 2 \cdot \begin{cases} q \cdot \left(y_{i,h} - \hat{y}_{i,h}^q\right), & \text{if } y_{i,h} < \hat{y}_{i,h}^q, \\ (1-q) \cdot \left(\hat{y}_{i,h}^q - y_{i,h}\right), & \text{if } y_{i,h} \geq \hat{y}_{i,h}^q, \end{cases} \tag{12}$$

and where a_i is the *historic absolute seasonal error* of the time series, defined as

$$a_i = \frac{1}{T-m} \sum_{t=m+1}^{T} \left|y_{i,t} - y_{i,t-m}\right|, \tag{13}$$

with m the seasonality of the dataset provided in Tab. 5.

- **Mean absolute scaled error (MASE)**:

$$\mathrm{MASE}(\hat{y}_i, y_i) \coloneqq \frac{1}{H} \frac{1}{a_i} \sum_{h=1}^{H} \left|\hat{y}_{i,T+h} - y_{i,T+h}\right|, \tag{14}$$

where a_i is the historic absolute seasonal error, as defined above. Note that MASE is equivalent to an SQL applied only to the 0.5 quantile.

To evaluate a particular time series model on a whole dataset, we average the loss across all individual items.

When evaluating point and probabilistic forecast performance, we used MASE or SQL as the training objective, respectively. This means, that all L2 and L3 models were trained once for point forecasting tasks using MASE as the training loss, and one more time for probabilistic forecasting tasks using the SQL loss.

C.2 Result aggregation

To summarize the empirical evaluation, we aggregate the individual losses $L_{m,d}$ of each model m and dataset d in a number of ways:

- **Elo**: The Elo rating system was originally introduced to calculate the relative skill levels of players in zero-sum games such as chess (Elo, 1967), but it has recently also been used to evaluate large language models (Bai et al., 2022; Boubdir et al., 2024). We compute our Elo scores as done in the recently proposed "Chatbot Arena" (Chiang et al., 2024), and calibrate the computation such that our chosen baseline has an Elo of 1000.

- **Average rank**: First, for each dataset in our evaluation, we rank the methods by their achieved error (resolving ties such that models receive the average rank of the tied models). Then, we compute the average rank of each model over all datasets with an arithmetic mean.

- **Champion**: This metric counts for how many datasets the method has achieved the lowest error among all methods included in the comparison. As we already computed a ranking for each dataset, this can be done by simply counting how often a method achieves rank 1. Since sometimes multiple methods are tied for the first place, the sum of the values in the "Champion" column does not always add up to the number of the datasets.

- **Average relative error**: To make the errors on different datasets more comparable, we first compute *relative* errors with respect to a chosen baseline method $\overline{m}$ (namely the "simple average (median)" method), as

$$L_{m,d}^{\text{rel}} \coloneqq \frac{L_{m,d}}{L_{\overline{m},d}}. \tag{15}$$

To limit the influence of outliers, we clip the individual relative errors to the range $[10^{-3}, 5]$. Then, we compute the *average relative error* of each model m across all datasets by aggregating the respective relative errors with a geometric mean:

$$\text{AverageRelativeError}_m = \text{GeometricMean}\left(L_{m,1}^{\text{rel}}, \ldots, L_{m,D}^{\text{rel}}\right). \tag{16}$$

D Computational resources

All models were trained on cloud-hosted machines with 16 vCPUs and 64GB RAM. The base models were trained with a time limit of 30 minutes per window per model to avoid extremely long runtimes, though the vast majority of models complete training within 5 minutes per window. The inidividual L2 stacker models were trained as described in Sec. B.2, either for a fixed, pre-specified number of steps, or until convergence within a time limit of 10 minutes. See also Tab. 11 and Tab. 12 for the median training times of the base and stacker models.

E Additional figures and tables

- Tab. 8: Full probabilistic forecasting results. SQL error values for all datasets × all methods (11 base models and 33 combination methods).

- Tab. 9: Full point forecasting results. MASE error values for all datasets × all methods (11 base models and 33 combination methods).

- Fig. 5: Critical differences (CD) diagram for 8 representative combination methods (probabilistic forecasting / SQL).

- Fig. 6: Critical differences (CD) diagram for 8 representative combination methods (point forecasting / MASE).

- Fig. 7: Distribution of relative error scores for 8 representative combination methods (probabilistic forecasting / SQL).

- Fig. 8: Distribution of relative error scores for 8 representative combination methods (point forecasting / MASE).

- Tab. 11: Median training time per validation window for L1 models.

- Tab. 12: Median end-to-end training time for the ensemble models.

- Fig. 9: Weights assigned by the L3 ensemble selection algorithm to the constituent L2 stacker models (probabilistic forecasting / SQL).

- Fig. 10: Weights assigned by the L3 ensemble selection algorithm to the constituent L2 stacker models (point forecasting / MASE).

- Tab. 10: Normalized error for each representative combination method and dataset combination (probabilistic forecasting / SQL).

E.1 Critical differences diagrams

Figure 5: Critical differences (CD) diagram for 8 representative combination methods. Based on probabilistic forecasting tasks (SQL metric).

Figure 6: Critical differences (CD) diagram for 8 representative combination methods. Based on point forecasting tasks (MASE metric).

E.2 Distribution of relative errors

Figure 7: Distribution of the relative errors for each combination method (normalized by the performance of median aggregation). Based on probabilistic forecasting tasks (SQL metric).

Figure 8: Distribution of the relative errors for each combination method (normalized by the performance of median aggregation). Based on point forecasting tasks (MASE metric).

Table 8: Quantile forecast accuracy for base and stacker models, measured with scaled quantile loss (SQL). Lower values are better.

Dataset	Individual models											Simple avg.		Model sel.	Performance-weighted			Greedy ensemble			Linear-softmax									Linear-positive									Nonlinear models				Stacker Model Selection	Multi-layer Stacking
	AutoETS	Theta	SeasonalNaive	PatchTST	DeepAR	DLinear	TiDE	DirectTabular	RecursiveTabular	TFT	Chronos-Bolt	Mean	Median	Model Selection	Exp.	Inv.	Sqr.	S=10	S=100	S=1000	m	mi	miq	miit	mitq	mq	mqq	mt	mtq	m	mi	miq	miit	mitq	mq	mqq	mt	mtq	RealMLP	RealMLP (scaled)	LightGBM	LightGBM (scaled)	Stacker Model Selection	Multi-layer Stacking
BEIJING-SUBWAY-30MIN																																												
Average rank																																												
Geometric relative error																																												

Datasets (rows): BEIJING-SUBWAY-30MIN, ETTh, ETTm, H2AIETRO, LOS-LOOP, M-DENSE, PEMS03, PEMS08, SHMETRO, SZ-TAXI, australian-electricity, bdg-2-bear, bdg-2-fox, bdg-2-panther, bdg-2-rat, beijing-air-quality, borealis, bull, cdc-fluview-ilinet, electricity-hourly, electricity-weekly, ercot, fred-md, gfc12-load, gfc17-load, hierarchical-sales, hog, ideal, kdd-cup-2018, kdd2022, m-daily, m4-daily, m4-hourly, m4-monthly, m4-quarterly, m4-weekly, m5, m5-weekly, pedestrian-counts, project-tycho, smart, solar-10-minutes, store-sales, subseasonal-precip, taxi-1h, taxi-30min, traffic-weekly, transactions, uber-tlc-daily, uber-tlc-hourly, wind-farms-daily.

Table 9: Point forecast accuracy for base and stacker models, measured with mean absolute scaled error (MASE). Lower values are better.

Table 10: Normalized error of each representative combination method per dataset, relative to the best model for that dataset (probabilstic tasks / SQL metric). Two key observations: (1) When most stackers perform reasonably (normalized error < 1.2), multi-layer stacking typically ranks among the top performers. (2) It underperforms when most constituent stackers—especially nonlinear ones—produce poor forecasts.

Dataset	Median	Model selection	Performance-based average	Greedy ensemble selection	Linear model	Nonlinear model	Stacker model selection	Multi-layer stacking
ETTh	1.068	1.173	1.049	1.056	1.050	1.157	1.025	1.000
HZMETRO	1.096	1.139	1.115	1.095	1.086	1.029	1.032	1.000
LOS-LOOP	1.364	1.301	1.146	1.104	1.095	1.046	1.056	1.000
PEMS08	1.378	1.124	1.028	1.056	1.048	1.039	1.039	1.000
bdg-2-bear	1.089	1.107	1.064	1.087	1.123	1.070	1.045	1.000
bdg-2-rat	1.240	1.051	1.067	1.051	1.038	1.055	1.016	1.000
borealis	1.090	1.068	1.051	1.047	1.055	1.014	1.017	1.000
hog	1.199	1.224	1.146	1.160	1.120	1.040	1.040	1.000
ideal	1.041	1.021	1.007	1.001	1.001	1.006	1.010	1.000
m4-daily	1.007	1.042	1.029	1.029	1.005	1.479	1.007	1.000
m4-hourly	1.350	1.149	1.149	1.101	1.089	1.029	1.027	1.000
m4-monthly	1.043	1.436	1.097	1.436	1.078	2.118	1.024	1.000
m4-quarterly	1.033	1.066	1.029	1.017	1.019	1.001	1.012	1.000
m4-weekly	1.173	1.250	1.098	1.084	1.073	1.109	1.042	1.000
nn5	1.057	1.025	1.008	1.007	1.008	1.021	1.008	1.000
nn5-weekly	1.043	1.046	1.017	1.025	1.026	1.083	1.025	1.000
project-tycho	1.020	1.011	1.025	1.004	1.001	1.003	1.001	1.000
store-sales	1.022	1.000	1.013	1.001	1.007	1.012	1.007	1.000
transactions	1.124	1.103	1.063	1.049	1.048	1.047	1.015	1.000
uber-tlc-daily	1.027	1.024	1.002	1.001	1.009	1.026	1.027	1.000
SZ-TAXI	1.027	1.055	1.014	1.001	1.000	1.011	1.001	1.001
PEMS03	1.076	1.000	1.067	1.064	1.049	1.108	1.076	1.005
kdd-cup-2018	1.163	1.046	1.026	1.002	1.000	1.092	1.000	1.005
hierarchical-sales	1.011	1.009	1.013	1.003	1.000	1.037	1.000	1.005
taxi-30min	1.188	1.002	1.003	1.002	1.000	1.008	1.008	1.005
bull	1.008	1.031	1.000	1.017	1.005	1.048	1.017	1.008
uber-tlc-hourly	1.151	1.000	1.003	1.000	1.001	1.016	1.013	1.010
pedestrian-counts	1.026	1.060	1.060	1.052	1.039	1.000	1.028	1.011
ercot	1.121	1.124	1.052	1.000	1.069	1.030	1.000	1.016
traffic-weekly	1.043	1.084	1.020	1.000	1.004	1.064	1.045	1.025
taxi-1h	1.479	1.000	1.000	1.006	1.003	1.046	1.006	1.026
smart	1.030	1.054	1.000	1.032	1.032	1.121	1.032	1.027
ETTm	1.124	1.000	1.052	1.037	1.026	1.230	1.042	1.027
electricity-weekly	1.290	1.000	1.050	1.017	1.041	1.175	1.017	1.030
beijing-air-quality	1.343	1.042	1.048	1.089	1.061	1.093	1.000	1.030
wind-farms-daily	1.030	1.000	1.040	1.005	1.011	1.074	1.074	1.036
subseasonal-precip	1.151	1.127	1.126	1.078	1.058	1.000	1.053	1.037
electricity-hourly	1.245	1.000	1.043	1.040	1.046	1.090	1.040	1.037
cdc-fluview-ilinet	1.072	1.077	1.000	1.015	1.027	1.201	1.110	1.048
BEIJING-SUBWAY-30MIN	1.055	1.607	1.466	1.072	1.073	1.000	1.090	1.051
bdg-2-panther	1.173	1.000	1.129	1.069	1.082	1.157	1.144	1.064
fred-md	1.000	1.127	1.011	1.015	1.000	1.009	1.088	1.068
SHMETRO	1.000	1.018	1.017	1.022	1.015	1.094	1.094	1.075
kdd2022	1.034	1.382	1.226	1.341	1.120	1.000	1.140	1.140
gfc12-load	1.077	1.299	1.023	1.000	1.060	1.361	1.133	1.140
bdg-2-fox	1.555	1.000	1.335	1.152	1.150	1.111	1.152	1.150
M-DENSE	1.058	1.052	1.009	1.000	1.071	1.270	1.270	1.159
australian-electricity	1.076	1.000	1.141	1.264	1.218	1.484	1.277	1.321
gfc17-load	1.000	2.568	1.179	1.317	1.464	1.385	1.991	1.447
solar-10-minutes	1.000	1.583	1.576	1.096	1.163	2.945	2.945	2.158

Table 11: Median training time for L1 models per one validation window (in seconds) for probabilistic tasks.

Model	Duration (seconds)
AutoETS	14
Chronos	8
DLinear	26
DeepAR	149
DirectTabular	233
Theta	8
PatchTST	63
RecursiveTabular	51
SeasonalNaive	1
TFT	204
TiDE	176

Table 12: Median end-to-end training time for ensemble models (in seconds) for probabilistic tasks.

Model	Duration (seconds)
Simple Average (Mean)	1
Simple Average (Median)	1
Model Selection	1
Performance-based (Exp.)	2
Performance-based (Inv.)	2
Performance-based (Sqr.)	2
Greed (S=10)	2
Greedy (S=100)	11
Greed (S=1000)	70
Linear (m, softmax)	15
Linear (m, positive)	13
Linear (mi, softmax)	7
Linear (mi, positive)	34
Linear (miq, softmax)	12
Linear (miq, positive)	44
Linear (mit, softmax)	12
Linear (mit, positive)	44
Linear (mitq, softmax)	30
Linear (mitq, positive)	77
Linear (mq, softmax)	11
Linear (mq, positive)	15
Linear (mqq, softmax)	10
Linear (mqq, positive)	33
Linear (mt, softmax)	7
Linear (mt, positive)	31
Linear (mtq, softmax)	13
Linear (mtq, positive)	52
LightGBM	25
LightGBM (scaled)	26
RealMLP	200
RealMLP (scaled)	207
Stacker Model Selection	484
Multi-layer Stacking	721

Heatmap of weights assigned by the L3 ensemble selection algorithm to the L2 stacker models across datasets.

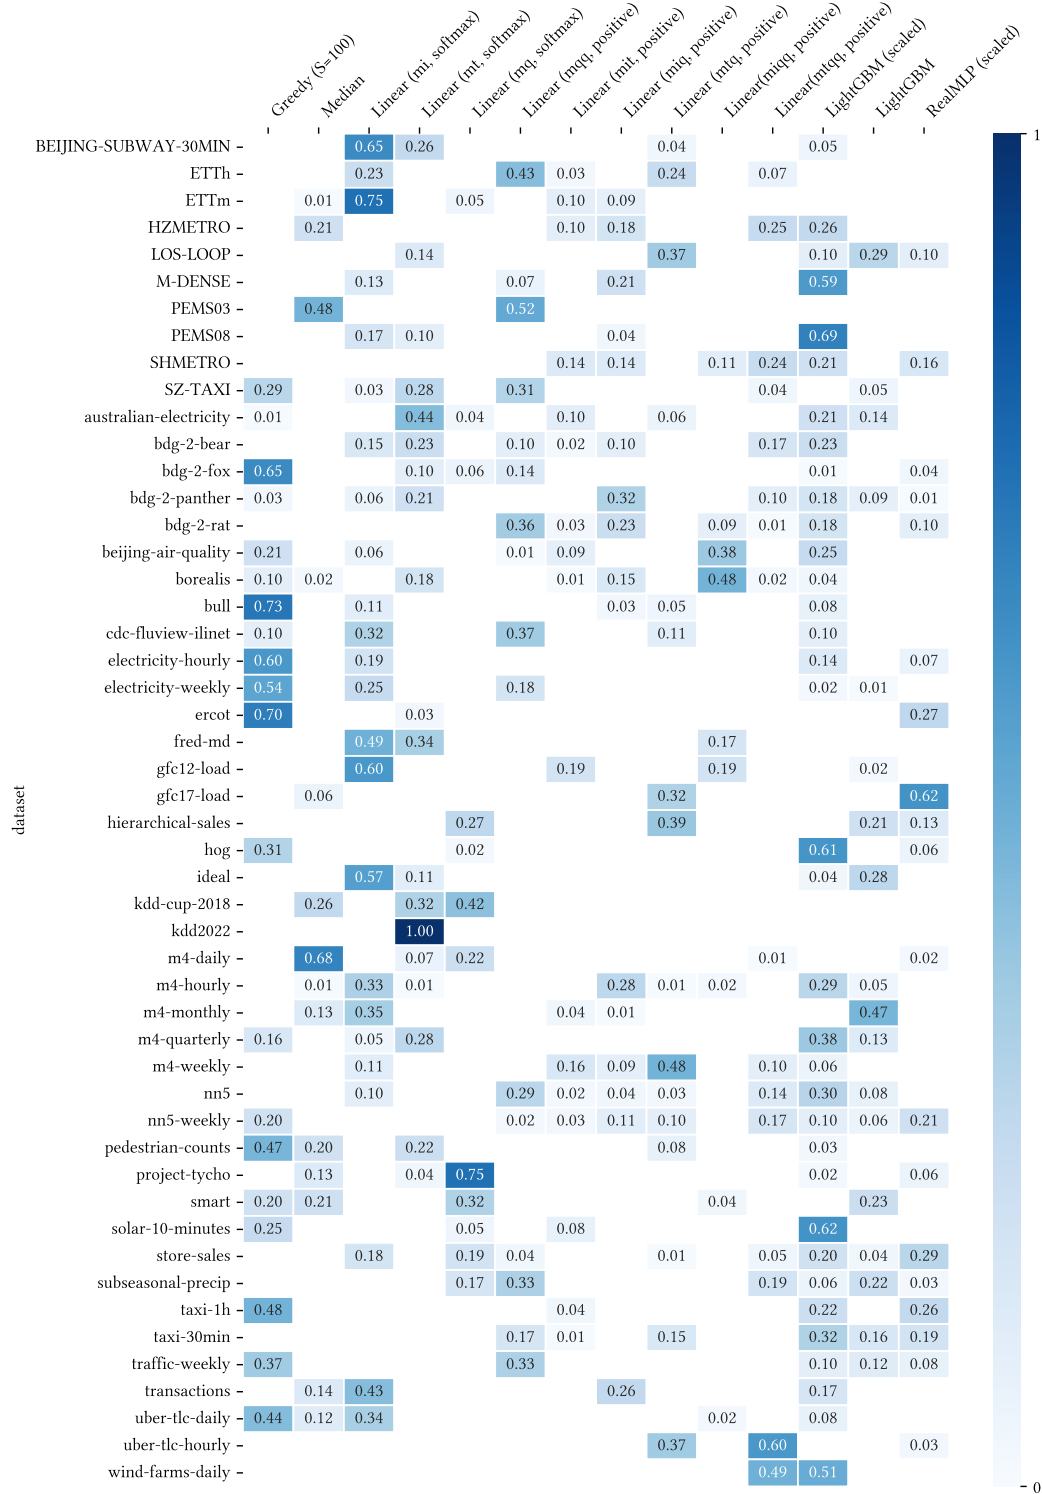| dataset	Greedy (S=100)	Median	Linear (mi, softmax)	Linear (mt, softmax)	Linear (mq, softmax)	Linear (mqq, positive)	Linear (mit, positive)	Linear (miq, positive)	Linear (mtq, positive)	Linear(miqq, positive)	Linear(mtqq, positive)	LightGBM (scaled)	LightGBM	RealMLP (scaled)
BEIJING-SUBWAY-30MIN			0.65	0.26				0.04				0.05		
ETTh			0.23			0.43	0.03	0.24		0.07				
ETTm	0.01		0.75		0.05		0.10	0.09						
HZMETRO		0.21					0.10	0.18			0.25	0.26		
LOS-LOOP				0.14					0.37			0.10	0.29	0.10
M-DENSE			0.13			0.07		0.21				0.59		
PEMS03	0.48					0.52								
PEMS08			0.17	0.10				0.04				0.69		
SHMETRO							0.14	0.14		0.11	0.24	0.21		0.16
SZ-TAXI	0.29		0.03	0.28		0.31					0.04		0.05	
australian-electricity	0.01			0.44	0.04		0.10		0.06			0.21	0.14	
bdg-2-bear			0.15	0.23		0.10	0.02	0.10			0.17	0.23		
bdg-2-fox	0.65			0.10	0.06	0.14						0.01		0.04
bdg-2-panther	0.03		0.06	0.21				0.32			0.10	0.18	0.09	0.01
bdg-2-rat						0.36	0.03	0.23		0.09	0.01	0.18		0.10
beijing-air-quality	0.21		0.06				0.01	0.09		0.38		0.25		
borealis	0.10	0.02		0.18			0.01	0.15		0.48	0.02	0.04		
bull	0.73		0.11					0.03	0.05			0.08		
cdc-fluview-ilinet	0.10		0.32			0.37			0.11			0.10		
electricity-hourly	0.60		0.19									0.14		0.07
electricity-weekly	0.54		0.25			0.18						0.02	0.01	
ercot	0.70			0.03										0.27
fred-md			0.49	0.34						0.17				
gfc12-load			0.60				0.19			0.19			0.02	
gfc17-load		0.06							0.32					0.62
hierarchical-sales					0.27				0.39				0.21	0.13
hog	0.31				0.02							0.61		0.06
ideal			0.57	0.11								0.04	0.28	
kdd-cup-2018		0.26		0.32	0.42									
kdd2022				1.00										
m4-daily		0.68		0.07	0.22						0.01			0.02
m4-hourly		0.01	0.33	0.01				0.28	0.01	0.02		0.29	0.05	
m4-monthly		0.13	0.35				0.04	0.01					0.47	
m4-quarterly	0.16		0.05	0.28								0.38	0.13	
m4-weekly			0.11				0.16	0.09	0.48		0.10	0.06		
nn5			0.10		0.29		0.02	0.04	0.03		0.14	0.30	0.08	
nn5-weekly	0.20					0.02	0.03	0.11	0.10		0.17	0.10	0.06	0.21
pedestrian-counts	0.47	0.20		0.22					0.08			0.03		
project-tycho		0.13		0.04	0.75							0.02		0.06
smart	0.20	0.21			0.32					0.04			0.23	
solar-10-minutes	0.25				0.05		0.08					0.62		
store-sales			0.18		0.19	0.04			0.01		0.05	0.20	0.04	0.29
subseasonal-precip					0.17	0.33					0.19	0.06	0.22	0.03
taxi-1h	0.48						0.04					0.22		0.26
taxi-30min							0.17	0.01	0.15			0.32	0.16	0.19
traffic-weekly	0.37				0.33							0.10	0.12	0.08
transactions		0.14	0.43					0.26				0.17		
uber-tlc-daily	0.44	0.12	0.34							0.02		0.08		
uber-tlc-hourly									0.37			0.60		0.03
wind-farms-daily											0.49	0.51		

Figure 9: Weights assigned by the L3 ensemble selection algorithm to the L2 stacker models over 50 probabilistic forecasting tasks.

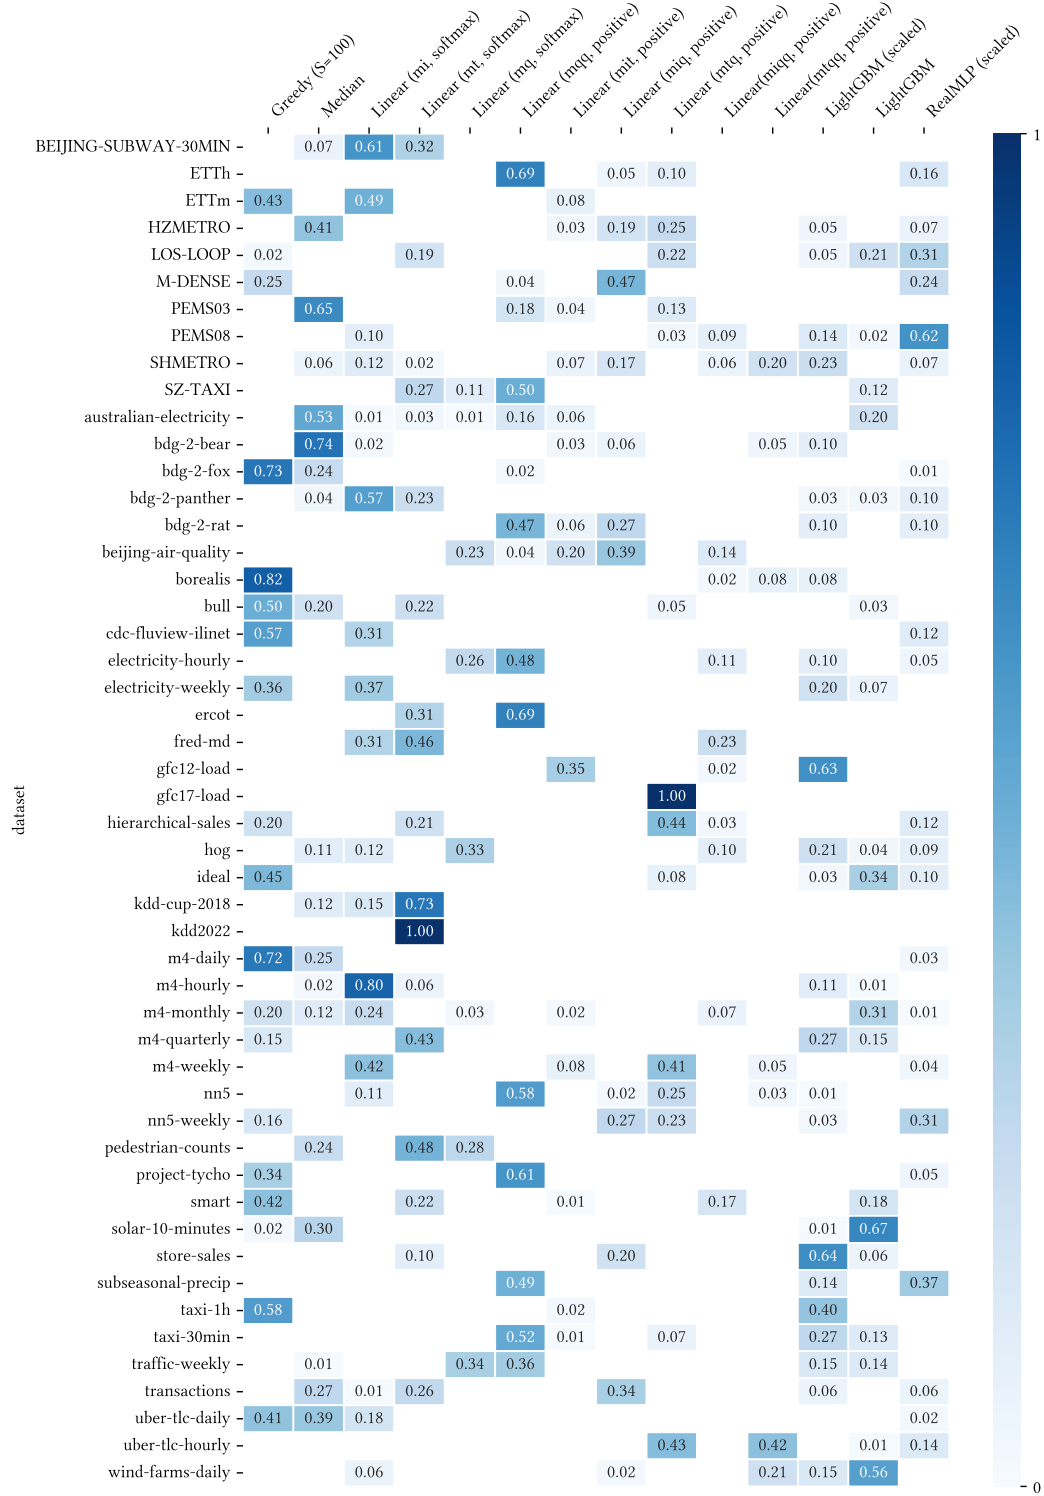| dataset	Greedy (S=100)	Median	Linear (mi, softmax)	Linear (mt, softmax)	Linear (mq, softmax)	Linear (mqq, positive)	Linear (mit, positive)	Linear (miq, positive)	Linear (mtq, positive)	Linear(miqq, positive)	Linear(mtqq, positive)	LightGBM (scaled)	LightGBM	RealMLP (scaled)
BEIJING-SUBWAY-30MIN		0.07	0.61	0.32										
ETTh					0.69		0.05	0.10						0.16
ETTm	0.43		0.49			0.08								
HZMETRO		0.41				0.03	0.19	0.25				0.05		0.07
LOS-LOOP	0.02			0.19				0.22				0.05	0.21	0.31
M-DENSE	0.25					0.04			0.47					0.24
PEMS03		0.65				0.18	0.04		0.13					0.62
PEMS08			0.10					0.03	0.09			0.14	0.02	0.62
SHMETRO		0.06	0.12	0.02			0.07	0.17		0.06	0.20	0.23		0.07
SZ-TAXI				0.27	0.11	0.50							0.12	
australian-electricity		0.53	0.01	0.03	0.01	0.16	0.06						0.20	
bdg-2-bear		0.74	0.02				0.03	0.06			0.05	0.10		
bdg-2-fox	0.73	0.24				0.02								0.01
bdg-2-panther		0.04	0.57	0.23								0.03	0.03	0.10
bdg-2-rat						0.47	0.06	0.27				0.10		0.10
beijing-air-quality					0.23	0.04	0.20	0.39		0.14				
borealis	0.82							0.02	0.08	0.08				
bull	0.50	0.20		0.22				0.05					0.03	
cdc-fluview-ilinet	0.57		0.31											0.12
electricity-hourly					0.26	0.48			0.11			0.10		0.05
electricity-weekly	0.36		0.37									0.20	0.07	
ercot				0.31		0.69								
fred-md			0.31	0.46					0.23					
gfc12-load							0.35		0.02		0.63			
gfc17-load								1.00						
hierarchical-sales	0.20			0.21				0.44	0.03					0.12
hog		0.11	0.12		0.33				0.10			0.21	0.04	0.09
ideal	0.45								0.08			0.03	0.34	0.10
kdd-cup-2018		0.12	0.15	0.73										
kdd2022				1.00										
m4-daily	0.72	0.25												0.03
m4-hourly		0.02	0.80	0.06								0.11	0.01	
m4-monthly	0.20	0.12	0.24		0.03		0.02		0.07				0.31	0.01
m4-quarterly	0.15			0.43								0.27	0.15	
m4-weekly			0.42				0.08		0.41		0.05			0.04
nn5		0.11			0.58		0.02	0.25			0.03	0.01		
nn5-weekly	0.16							0.27	0.23			0.03		0.31
pedestrian-counts		0.24		0.48	0.28									
project-tycho	0.34					0.61								0.05
smart	0.42			0.22			0.01			0.17		0.18		
solar-10-minutes	0.02	0.30										0.01	0.67	
store-sales				0.10				0.20				0.64	0.06	
subseasonal-precip						0.49						0.14		0.37
taxi-1h	0.58						0.02					0.40		
taxi-30min						0.52	0.01		0.07			0.27	0.13	
traffic-weekly		0.01			0.34	0.36						0.15	0.14	
transactions		0.27	0.01	0.26					0.34			0.06		0.06
uber-tlc-daily	0.41	0.39	0.18											0.02
uber-tlc-hourly									0.43		0.42		0.01	0.14
wind-farms-daily			0.06				0.02				0.21	0.15	0.56	

Figure 10: Weights assigned by the L3 ensemble selection algorithm to the L2 stacker models over 50 point forecasting tasks.